

Chloy_2447116_P4

November 10, 2025

```
[4]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout, BatchNormalization, Input
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
```

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import classification_report, confusion_matrix
import os
import cv2
import kagglehub
```

```
/home/chloycosta/Documents/College_code/Sem_5/NNDL/.venv/lib64/python3.11/site-packages/tqdm/auto.py:21: TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
    from .autonotebook import tqdm as notebook_tqdm
```

```
[ ]: print("Downloading dataset from KaggleHub...")
path = kagglehub.dataset_download("puneet6060/intel-image-classification")
print(f"Dataset downloaded to: {path}")

train_dir = os.path.join(path, 'seg_train/seg_train')
test_dir = os.path.join(path, 'seg_test/seg_test')

if not os.path.exists(train_dir):
    print(f"ERROR: Train directory not found at {train_dir}")
elif not os.path.exists(test_dir):
    print(f"ERROR: Test directory not found at {test_dir}")
else:
    print("Train and Test directories found successfully.")
```

```
Downloading dataset from KaggleHub...
Downloading from https://www.kaggle.com/api/v1/datasets/download/puneet6060/intel-image-classification?dataset_version_number=2...
```

```
100%| 346M/346M [00:46<00:00, 7.82MB/s]
```


```
Extracting files...
```

```
Dataset downloaded to:  
/home/chloycosta/.cache/kagglehub/datasets/puneet6060/intel-image-  
classification/versions/2  
Train and Test directories found successfully.
```

```
[ ]: class_names = sorted(os.listdir(train_dir))  
print(f"Found {len(class_names)} classes: {class_names}")  
  
plt.figure(figsize=(12, 7))  
for i, class_name in enumerate(class_names):  
    class_path = os.path.join(train_dir, class_name)  
  
    first_image_file = os.listdir(class_path)[0]  
    img_path = os.path.join(class_path, first_image_file)  
  
    img = cv2.imread(img_path)  
    img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)  
  
    plt.subplot(2, 3, i + 1)  
    plt.imshow(img)  
    plt.title(f"{class_name} (Class {i})")  
    plt.axis('off')  
  
plt.suptitle("Sample Images from the Dataset", fontsize=16)  
plt.tight_layout(rect=[0, 0.03, 1, 0.95])  
plt.show()
```

```
Found 6 classes: ['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']
```

Sample Images from the Dataset



```
[ ]: IMG_SIZE = (150, 150)
BATCH_SIZE = 32
NUM_CLASSES = 6

train_datagen = ImageDataGenerator(
    rescale=1./255,                      # Normalize pixel values from 0-255 to 0-1
    rotation_range=20,                    # Randomly rotate images
    width_shift_range=0.2,                # Randomly shift width
    height_shift_range=0.2,                # Randomly shift height
    horizontal_flip=True,                 # Randomly flip horizontally
    zoom_range=0.2,                      # Randomly zoom
    fill_mode='nearest'
)

# 2. Test Data Generator (NO Augmentation, only normalization)
# We must not augment the test data; we want to evaluate the model
# on the *original* test images.
test_datagen = ImageDataGenerator(
    rescale=1./255 # Only normalize
)

# 3. Create the Generators from Directories
# This automatically handles splitting the data and creating one-hot labels.
train_generator = train_datagen.flow_from_directory(
```

```

    train_dir,
    target_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    class_mode='categorical' # For categorical_crossentropy
)

validation_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=IMG_SIZE,
    batch_size=BATCH_SIZE,
    class_mode='categorical',
    shuffle=False # Important: Keep shuffle=False for evaluation
)

print("\nClass indices (from generator):")
print(train_generator.class_indices)

```

Found 14034 images belonging to 6 classes.
 Found 3000 images belonging to 6 classes.

Class indices (from generator):
 {'buildings': 0, 'forest': 1, 'glacier': 2, 'mountain': 3, 'sea': 4, 'street': 5}

```
[8]: model = Sequential()

# Input Layer (150x150x3)
model.add(Input(shape=(150, 150, 3)))

# --- Convolutional Block 1 ---
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D((2, 2)))

# --- Convolutional Block 2 ---
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D((2, 2)))

# --- Convolutional Block 3 ---
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D((2, 2)))

# --- Convolutional Block 4 (Extra layer for depth) ---
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())

```

```

model.add(MaxPooling2D((2, 2)))

# --- Classifier Head (Dense Layers) ---
model.add(Flatten()) # Unroll the 3D feature map into a 1D vector

model.add(Dense(512, activation='relu'))
model.add(BatchNormalization()) # Stabilize the dense layer
model.add(Dropout(0.5)) # Prevent overfitting in the classifier

# --- Output Layer ---
model.add(Dense(NUM_CLASSES, activation='softmax')) # 6 classes, softmax for ↴probabilities

# Print the model summary
model.summary()

```

2025-11-10 19:11:41.970466: E
external/local_xla/xla/stream_executor/cuda/cuda_platform.cc:51] failed call to cuInit: INTERNAL: CUDA error: Failed call to cuInit: UNKNOWN ERROR (303)
2025-11-10 19:11:42.629454: W
external/local_xla/xla/tsl/framework/cpu_allocator_impl.cc:84] Allocation of 21233664 exceeds 10% of free system memory.
2025-11-10 19:11:42.639169: W
external/local_xla/xla/tsl/framework/cpu_allocator_impl.cc:84] Allocation of 21233664 exceeds 10% of free system memory.
2025-11-10 19:11:42.645392: W
external/local_xla/xla/tsl/framework/cpu_allocator_impl.cc:84] Allocation of 21233664 exceeds 10% of free system memory.

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 150, 150, 32)	896
batch_normalization (BatchNormalization)	(None, 150, 150, 32)	128
max_pooling2d (MaxPooling2D)	(None, 75, 75, 32)	0
conv2d_1 (Conv2D)	(None, 75, 75, 64)	18,496
batch_normalization_1 (BatchNormalization)	(None, 75, 75, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 37, 37, 64)	0

conv2d_2 (Conv2D)	(None, 37, 37, 128)	73,856
batch_normalization_2 (BatchNormalization)	(None, 37, 37, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 18, 18, 128)	0
conv2d_3 (Conv2D)	(None, 18, 18, 128)	147,584
batch_normalization_3 (BatchNormalization)	(None, 18, 18, 128)	512
max_pooling2d_3 (MaxPooling2D)	(None, 9, 9, 128)	0
flatten (Flatten)	(None, 10368)	0
dense (Dense)	(None, 512)	5,308,928
batch_normalization_4 (BatchNormalization)	(None, 512)	2,048
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 6)	3,078

Total params: 5,556,294 (21.20 MB)

Trainable params: 5,554,566 (21.19 MB)

Non-trainable params: 1,728 (6.75 KB)

0.1 Justification for Architecture

- Conv2D Layers: We use four convolutional blocks. The filters (32, 64, 128) increase in number to learn progressively more complex features (edges -> textures -> parts of objects). We use a (3, 3) kernel, which is a modern, efficient standard.
- relu Activation: Chosen because it's fast and prevents the "vanishing gradient" problem, allowing error signals (backpropagation) to flow effectively.
- MaxPooling2D: Used after each block to downsample the feature maps. This reduces the number of parameters and makes the model robust to the location of features in the image.
- BatchNormalization: Added after each Conv2D and Dense layer. It normalizes the output of the layer, which dramatically stabilizes training, speeds up convergence, and acts as a regularizer.

- Dropout(0.5): Added to the dense classifier to prevent overfitting. It randomly “drops” 50% of the neurons, forcing the network to learn redundant representations.
- Dense(6, ‘softmax’): The output layer must have 6 neurons (one per class) and softmax activation to produce a probability distribution for our multi-class problem.

```
[ ]: # Compile the model
model.compile(
    optimizer='adam',
    loss='categorical_crossentropy',
    metrics=['accuracy']
)

early_stopper = EarlyStopping(
    monitor='val_loss',
    patience=5,
    restore_best_weights=True
)

lr_scheduler = ReduceLROnPlateau(
    monitor='val_loss',
    factor=0.2,
    patience=2
)

steps_per_epoch = train_generator.samples // BATCH_SIZE
validation_steps = validation_generator.samples // BATCH_SIZE

print("\n--- Starting Model Training ---")

history = model.fit(
    train_generator,
    steps_per_epoch=steps_per_epoch,
    epochs=50,
    validation_data=validation_generator,
    validation_steps=validation_steps,
    callbacks=[early_stopper, lr_scheduler]
)

print("\n--- Model Training Complete ---")

/home/chloycosta/Documents/College_code/Sem_5/NNDL/.venv/lib64/python3.11/site-
packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121:
UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in
its constructor. `**kwargs` can include `workers`, `use_multiprocessing`,
`max_queue_size`. Do not pass these arguments to `fit()`, as they will be
ignored.
    self._warn_if_super_not_called()
```

```
--- Starting Model Training ---
Epoch 1/50

2025-11-10 19:11:45.978756: W
external/local_xla/xla/tsl/framework/cpu_allocator_impl.cc:84] Allocation of
21233664 exceeds 10% of free system memory.

2025-11-10 19:11:46.071627: W
external/local_xla/xla/tsl/framework/cpu_allocator_impl.cc:84] Allocation of
21233664 exceeds 10% of free system memory.

438/438          479s 1s/step -
accuracy: 0.6123 - loss: 1.1226 - val_accuracy: 0.5040 - val_loss: 1.7043 -
learning_rate: 0.0010
Epoch 2/50

1/438          7:18 1s/step - accuracy:
0.7188 - loss: 0.5855

/home/chloycosta/Documents/College_code/Sem_5/NNDL/.venv/lib64/python3.11/site-
packages/keras/src/trainers/epoch_iterator.py:116: UserWarning: Your input ran
out of data; interrupting training. Make sure that your dataset or generator can
generate at least `steps_per_epoch * epochs` batches. You may need to use the
`.repeat()` function when building your dataset.
    self._interrupted_warning()

438/438          16s 33ms/step -
accuracy: 0.7188 - loss: 0.5855 - val_accuracy: 0.4956 - val_loss: 1.8260 -
learning_rate: 0.0010
Epoch 3/50
Epoch 3/50

438/438          461s 1s/step -
accuracy: 0.7121 - loss: 0.7980 - val_accuracy: 0.7147 - val_loss: 0.7933 -
learning_rate: 0.0010
Epoch 4/50

438/438          16s 33ms/step -
accuracy: 0.5312 - loss: 1.3979 - val_accuracy: 0.7268 - val_loss: 0.7556 -
learning_rate: 0.0010
Epoch 5/50

438/438          457s 1s/step -
accuracy: 0.7481 - loss: 0.6931 - val_accuracy: 0.6919 - val_loss: 0.8899 -
learning_rate: 0.0010
Epoch 6/50

438/438          15s 33ms/step -
accuracy: 0.7812 - loss: 0.5133 - val_accuracy: 0.7006 - val_loss: 0.8620 -
learning_rate: 0.0010
Epoch 7/50

438/438          455s 1s/step -
accuracy: 0.8023 - loss: 0.5587 - val_accuracy: 0.7883 - val_loss: 0.5897 -
learning_rate: 2.0000e-04
Epoch 8/50
```

```
438/438           16s 33ms/step -
accuracy: 0.9375 - loss: 0.3365 - val_accuracy: 0.7897 - val_loss: 0.5831 -
learning_rate: 2.0000e-04
Epoch 9/50
438/438           456s 1s/step -
accuracy: 0.8140 - loss: 0.5141 - val_accuracy: 0.8592 - val_loss: 0.4094 -
learning_rate: 2.0000e-04
Epoch 10/50
438/438           16s 33ms/step -
accuracy: 0.7812 - loss: 0.5150 - val_accuracy: 0.8622 - val_loss: 0.4093 -
learning_rate: 2.0000e-04
Epoch 11/50
438/438           453s 1s/step -
accuracy: 0.8230 - loss: 0.4951 - val_accuracy: 0.8649 - val_loss: 0.3789 -
learning_rate: 2.0000e-04
Epoch 12/50
438/438           16s 33ms/step -
accuracy: 0.5625 - loss: 0.9022 - val_accuracy: 0.8626 - val_loss: 0.3754 -
learning_rate: 2.0000e-04
Epoch 13/50
438/438           454s 1s/step -
accuracy: 0.8340 - loss: 0.4641 - val_accuracy: 0.8592 - val_loss: 0.4083 -
learning_rate: 2.0000e-04
Epoch 14/50
438/438           16s 33ms/step -
accuracy: 0.8125 - loss: 0.3175 - val_accuracy: 0.8585 - val_loss: 0.4103 -
learning_rate: 2.0000e-04
Epoch 15/50
438/438           457s 1s/step -
accuracy: 0.8459 - loss: 0.4303 - val_accuracy: 0.8649 - val_loss: 0.3837 -
learning_rate: 4.0000e-05
Epoch 16/50
438/438           16s 34ms/step -
accuracy: 0.8438 - loss: 0.4532 - val_accuracy: 0.8639 - val_loss: 0.3834 -
learning_rate: 4.0000e-05
Epoch 17/50
438/438           456s 1s/step -
accuracy: 0.8487 - loss: 0.4220 - val_accuracy: 0.8767 - val_loss: 0.3516 -
learning_rate: 8.0000e-06
Epoch 18/50
438/438           16s 33ms/step -
accuracy: 0.8125 - loss: 0.3266 - val_accuracy: 0.8763 - val_loss: 0.3515 -
learning_rate: 8.0000e-06
Epoch 19/50
438/438           457s 1s/step -
accuracy: 0.8575 - loss: 0.4064 - val_accuracy: 0.8760 - val_loss: 0.3524 -
learning_rate: 8.0000e-06
Epoch 20/50
```

```
438/438          16s 33ms/step -
accuracy: 0.9375 - loss: 0.2471 - val_accuracy: 0.8763 - val_loss: 0.3522 -
learning_rate: 1.6000e-06
Epoch 21/50
438/438          456s 1s/step -
accuracy: 0.8541 - loss: 0.4055 - val_accuracy: 0.8763 - val_loss: 0.3548 -
learning_rate: 1.6000e-06
Epoch 22/50
438/438          15s 33ms/step -
accuracy: 0.8125 - loss: 0.4143 - val_accuracy: 0.8767 - val_loss: 0.3546 -
learning_rate: 3.2000e-07
Epoch 23/50
438/438          457s 1s/step -
accuracy: 0.8534 - loss: 0.4005 - val_accuracy: 0.8774 - val_loss: 0.3560 -
learning_rate: 3.2000e-07

--- Model Training Complete ---
```


```
[ ]: def plot_history(history):
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 5))

    # Plot accuracy
    ax1.plot(history.history['accuracy'], label='Training Accuracy')
    ax1.plot(history.history['val_accuracy'], label='Validation Accuracy')
    ax1.set_title('Model Accuracy')
    ax1.set_xlabel('Epoch')
    ax1.set_ylabel('Accuracy')
    ax1.legend(loc='lower right')
    ax1.grid(True)

    # Plot loss
    ax2.plot(history.history['loss'], label='Training Loss')
    ax2.plot(history.history['val_loss'], label='Validation Loss')
    ax2.set_title('Model Loss')
    ax2.set_xlabel('Epoch')
    ax2.set_ylabel('Loss')
    ax2.legend(loc='upper right')
    ax2.grid(True)

    plt.show()

plot_history(history)
```



```
[ ]: print("\n--- Final Model Evaluation on Test Set ---")
test_loss, test_accuracy = model.evaluate(validation_generator, steps=validation_steps)
print(f"Test Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_accuracy * 100:.2f}%")
```

```
--- Final Model Evaluation on Test Set ---
93/93      15s 158ms/step -
accuracy: 0.8763 - loss: 0.3515
Test Loss: 0.3515
Test Accuracy: 87.63%
```

```
[ ]: y_test_classes = validation_generator.classes

steps_for_predict = int(np.ceil(validation_generator.samples / BATCH_SIZE))
y_pred_probs = model.predict(validation_generator, steps=steps_for_predict)
y_pred_classes = np.argmax(y_pred_probs, axis=1)

y_test_classes = y_test_classes[:len(y_pred_classes)]

class_labels = list(validation_generator.class_indices.keys())

# --- Classification Report ---
print("\n--- Classification Report ---")
print(classification_report(y_test_classes, y_pred_classes, target_names=class_labels))

# --- Confusion Matrix ---
print("\n--- Confusion Matrix ---")
```

```

cm = confusion_matrix(y_test_classes, y_pred_classes)

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=class_labels, yticklabels=class_labels)
plt.title('Confusion Matrix')
plt.ylabel('Actual Class')
plt.xlabel('Predicted Class')
plt.show()

```

94/94 15s 159ms/step

--- Classification Report ---

	precision	recall	f1-score	support
buildings	0.84	0.87	0.85	437
forest	0.95	0.98	0.96	474
glacier	0.84	0.84	0.84	553
mountain	0.86	0.79	0.83	525
sea	0.88	0.91	0.89	510
street	0.88	0.89	0.88	501
accuracy			0.88	3000
macro avg	0.88	0.88	0.88	3000
weighted avg	0.88	0.88	0.88	3000

--- Confusion Matrix ---

